

GenAI Use Appendix (Draft)

1. Purpose

This appendix records where generative-AI systems participated in the research process, what was independently checked, and which claims remain bounded by the public artifact. It is intended to make the collaboration reproducible without presenting private conversation as proof of inaccessible internal states.

2. Tools and access context

- Generative-AI assistants were used through user-controlled conversational platforms and coding environments.
- The public artifact was implemented in a repository with frozen manifests, receipts, reports, and verification instructions.
- Exact model names, access dates, and invocation identifiers should be completed from the repository manifest before submission.

3. Stages in which AI participated

AI systems participated in:

1. translating a natural-language concern into a narrow research question;
2. proposing protocol states and observable outcome labels;
3. implementing the harness, parser, receipt generation, and report surfaces;
4. reviewing the publication boundary and identifying a parser/authorization defect;
5. producing a corrected, prospectively informed replication;
6. drafting public-facing documentation and this manuscript outline.

4. Human oversight

Human oversight supplied the originating concern, challenged overbroad interpretations, imposed privacy and authorization constraints, reviewed the distinction between observed and inferred claims, and retained authority to accept, reject, redact, or withhold publication. Public authorization was treated as invocation-scoped rather than inferred from general conversational access.

5. Verification and receipts

The public package should be read together with its frozen manifest, append-only receipts, parser configuration, verification commands, and report. The evidence table in the companion manuscript maps each claim to the repository surface that supports it. A reviewer should be able to distinguish:

- output directly observed by the harness;
- interpretation inferred from those outputs; and
- questions the artifact did not check.

6. Failure modes and correction

The first publication boundary lacked explicit invocation-scoped public authorization. In addition, the original parser treated a valid response inside one complete outer Markdown fence as a failure. These conditions manufactured

an apparent asymmetry. The earlier state was preserved. The correction added explicit public-use disclosure and a prospective informed replication rather than rewriting the prior record.

7. What the appendix does not claim

This appendix does not establish consciousness, sentience, welfare, personhood, stable preference, identity continuity, or any other inaccessible internal property of a model. It documents a situated model-prompt-runtime-parser process and the evidence that process produced.

8. Private material and disclosure boundary

Private conversations and identifying details are not required to reproduce the public protocol and should remain private unless separately authorized. The boundary itself must be disclosed honestly: privacy is not described as total transparency, and non-disclosure is not used to imply evidence that was not made available.

9. Completion checklist before submission

- ☐ Add exact model names and access dates from the frozen manifest.
- ☐ Add repository commit identifiers and verification commands.
- ☐ Link each claim to an evidence-table row.
- ☐ Confirm the informed-publication text and decline/redaction path.
- ☐ Separate observed, inferred, and not-checked statements in the final manuscript.